

10

Advanced SQL Query Optimization Techniques

Every Data Scientist Must Know!

*Shruthi
Annamaneni*



like



share



comment

Struggling with slow SQL queries?

As a data scientist, optimizing your database performance is crucial for efficiency and scalability. Here are 10 proven SQL query optimization techniques that will take your skills to the next level!

1

Index Optimization

- Ensure indexes are created on columns that are frequently used in '**WHERE**' clauses, '**JOIN**' conditions and as part of '**ORDER BY**' clauses.
- Use composite indexes for columns that are frequently queried together.
- Regularly analyze and rebuild fragmented indexes.

2

Query Refactoring

- Break complex queries into simpler subqueries or use common table expressions (CTEs).
- Avoid unnecessary columns in the 'SELECT' clause to reduce the data processed.

3

Join Optimization

- Use the appropriate type of join (INNER JOIN, LEFT JOIN, etc.) based on the requirements.
- Ensure join columns are indexed to speed up the join operation.
- Consider the join order, starting with the smallest table.

4

Use of Proper Data Types

- Choose the most efficient data type for your columns to reduce storage and improve performance.
- Avoid using 'SELECT *', specify only the columns you need.

5

Query Execution Plan Analysis

- Use tools like 'EXPLAIN' or 'EXPLAIN PLAN' to analyze how the database executes a query.
- Look for full table scans, inefficient joins, or unnecessary sorting operations.

6

Temporary Tables and Materialized Views

- Use temporary tables to store intermediate results that are reused multiple times in complex queries.
- Use materialized views to store precomputed results of expensive queries.

7

Efficient Use of Subqueries and CTEs

- Replace correlated subqueries with joins when possible to avoid repeated execution.
- Use CTEs to improve readability and reusability, and sometimes performance, of complex queries.

8

Optimization of Aggregate Functions

- Use indexed columns in 'GROUP BY' clauses to speed up aggregation.
- Consider using window functions for complex aggregations instead of traditional 'GROUP BY'.

9

Avoiding Functions in Predicates

- Avoid using functions on columns in the 'WHERE' clause as it can prevent the use of indexes.
- Rewrite conditions to allow the use of indexes.

10

Parameter Sniffing and Query Caching

- Be aware of parameter sniffing issues where SQL Server caches execution plans based on initial parameter values.
- Use query hints or option recompile to address specific performance issues.
- Take advantage of query caching mechanisms where appropriate to reuse execution plans.